




Three Agent loop strategies as rollout pipes

Each implementation has its own inner interaction pipe and outer runtime. The large return arrow is the rollout boundary from output at t to next input at $t + 1$.



nanoGPT placement: runner.ts currently spans the pipe plus session state/persistence and task verification. A cleaner target is inner AgentLoop → outer SessionRunner → harness verifier

Color is redundant with explicit state labels; arrows encode control flow, not performance or reward magnitude.